

Language Understanding for Text-based Games using Deep Reinforcement Learning

Karthik Narasimhan, Tejas D Kulkarni, Regina Barzilay

Presented by Mark Chang

Outline

- Introduction
- Background
 - Reinforcement Learning
 - Q-Learning
 - Deep Q-Network
 - Long Short-Term Memory
- Deep Reinforcement Learning Framework
 - Game Representation
 - Representation Generator
 - Action Scorer
 - Parameter Learning
- Experiment & Result
- Conclusion

Introduction

- Learning the control policies for text-based games.
- A deep **reinforcement learning** framework
 - learn state representations (semantics of text)
 - learn action policies (rules of game)

Introduction

- MUD(Multi-User Dungeon) Game
 - A text-based multiplayer real-time virtual world.
 - Player read descriptions of the world state.
 - Player take actions by natural language commands.

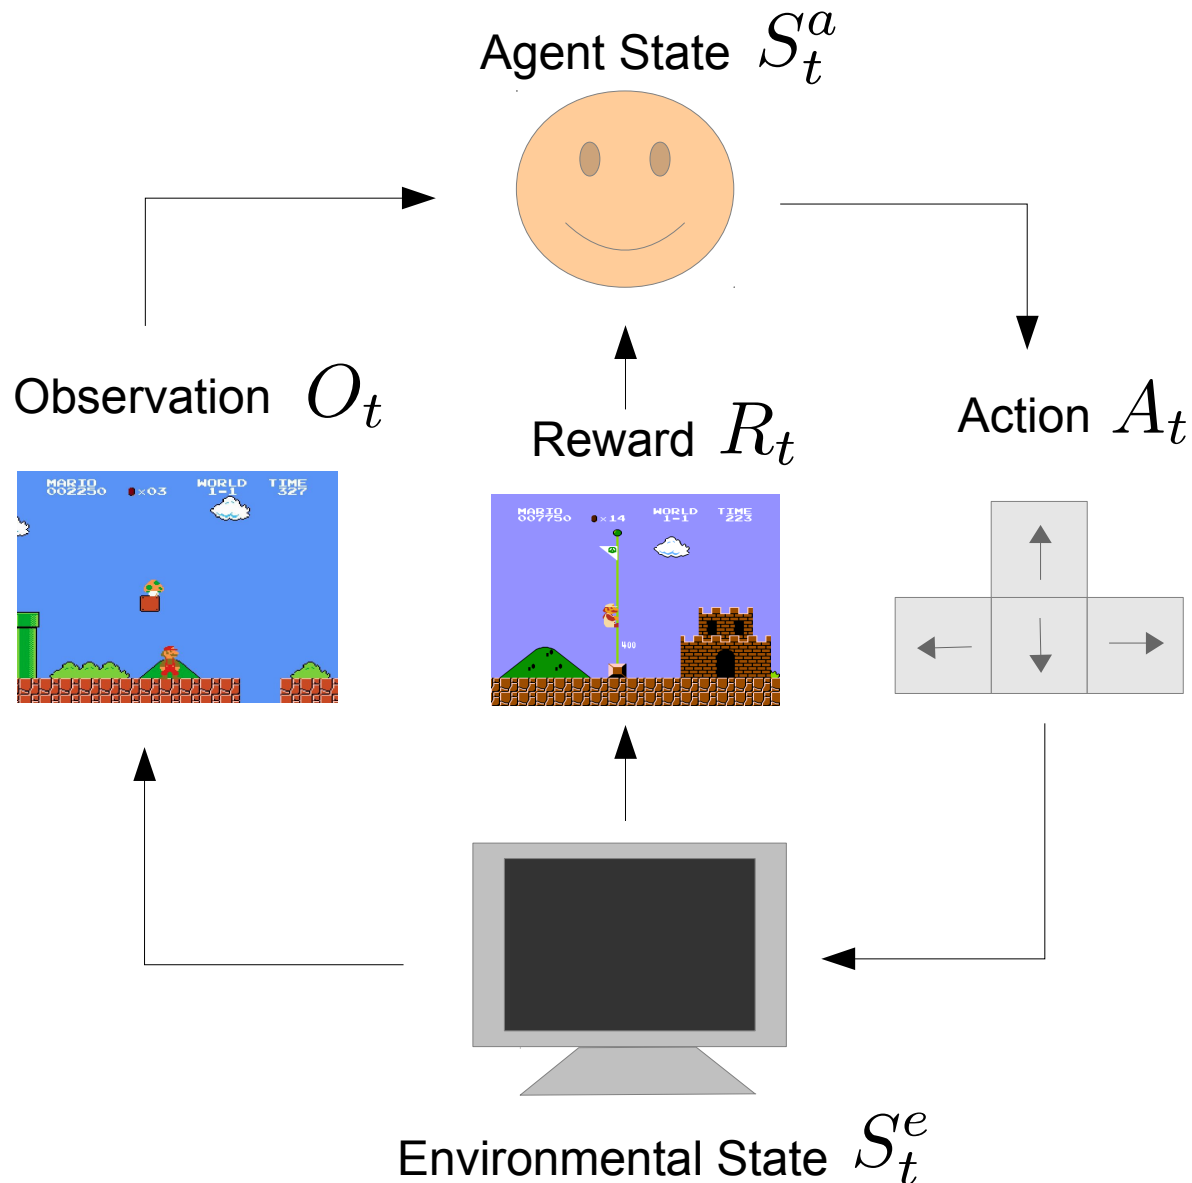
You are standing very close to the bridge's eastern foundation.
If you go east you will be back on solid ground.
The bridge sways in the wind.

> Go east

Background

- Reinforcement Learning
- Q-Learning
- Deep Q-Network
- Long Short-Term Memory

Reinforcement Learning



- Rules of the game are unknown.
- No supervisor, only a reward signal.
- Feedback is delayed.
- Agent's actions affect the subsequent data it receives.

Reinforcement Learning

- Markov Decision Process $\langle S, A, P, R, \gamma \rangle$
 - S is a finite set of states
 - A is a finite set of actions
 - P is a state transition probability matrix
$$P_{ss'}^a = P[S_{t+1} = s' \mid S_t = s, A_t = a]$$
 - R is a reward function
$$R_s = E[R_{t+1} \mid S_t = s]$$
 - γ is a discount factor


Reinforcement Learning

- Return : total discounted reward from time-step t .
 - $G_t = R_{t+1} + \gamma R_{t+2} + \dots + R_{\text{final}}$
- Policy : Agent Behavior
 - Deterministic policy: $\pi(s) = a$
 - Stochastic policy: $\pi(a \mid s) = P[A_t = a \mid S_t = s]$
- Value function: Expected return starting from state s .
 - state-value function: $v_\pi(s) = E_\pi[G_t \mid S_t = s]$
 - action-value function: $q_\pi(s, a) = E_\pi[G_t \mid S_t = s, A_t = a]$

Reinforcement Learning

- Bellman Equation for Action-Value Function

$$q_{\pi}(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a \sum_{a' \in A} \pi(a' | s') q_{\pi}(s', a')$$




Q-Learning

- A model-free, off-policy technique to learn an optimal $Q(s, a)$ for the agent.

$$Q_{i+1}(s, a) \leftarrow Q_i(s, a) + \alpha(R + \gamma \max_{a'} Q_i(s', a') - Q_i(s, a))$$

- model-free: using sampling instead of complete model
- off-policy: learning optimal policy independently of agent's actions
- $Q(s, a)$: Estimation of $q(s, a)$
- α : Step Size (Learning Rate)



Q-Learning

- How to choose A in $Q(s, a)$?
- ϵ - Greedy Exploration
 - ensuring continual exploration
 - With probability $1 - \epsilon$ choose the greedy action
 - With probability ϵ choose an action at random

$$\pi(a \mid s) = \begin{cases} \frac{\epsilon}{m} + 1 - \epsilon & \text{if } a = \operatorname{argmax}_{a'} Q(s, a') \\ \frac{\epsilon}{m} & \text{otherwise} \end{cases}$$

Deep Q-Network

- Problem with $Q(s, a)$:
 - too many states/actions to store in memory
 - too slow to individually learn their value
 - Solution :
 - Approximate $Q(s, a)$ using a neural network
- Q-Network : $Q(s, a; \theta)$

Deep Q-Network

- Loss Function:

reduce the discrepancy between

- the predicted value : $Q(s, a; \theta_i)$
- the expected Q-value : $r + \gamma \max_{a'} Q(s', a'; \theta_{i-1})$

$$L(\theta_i) = E_{s,a}[(r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i))^2]$$


- Gradient of Loss Function:

$$\nabla_{\theta_i} L(\theta_i) = E_{s,a}[(r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i)) \nabla_{\theta_i} Q(s, a; \theta_i)]$$

- previous iteration θ_{i-1} is fixed


Long Short-Term Memory

- recurrent neural networks with memory
- able to connect and recognize long-range patterns between words in text

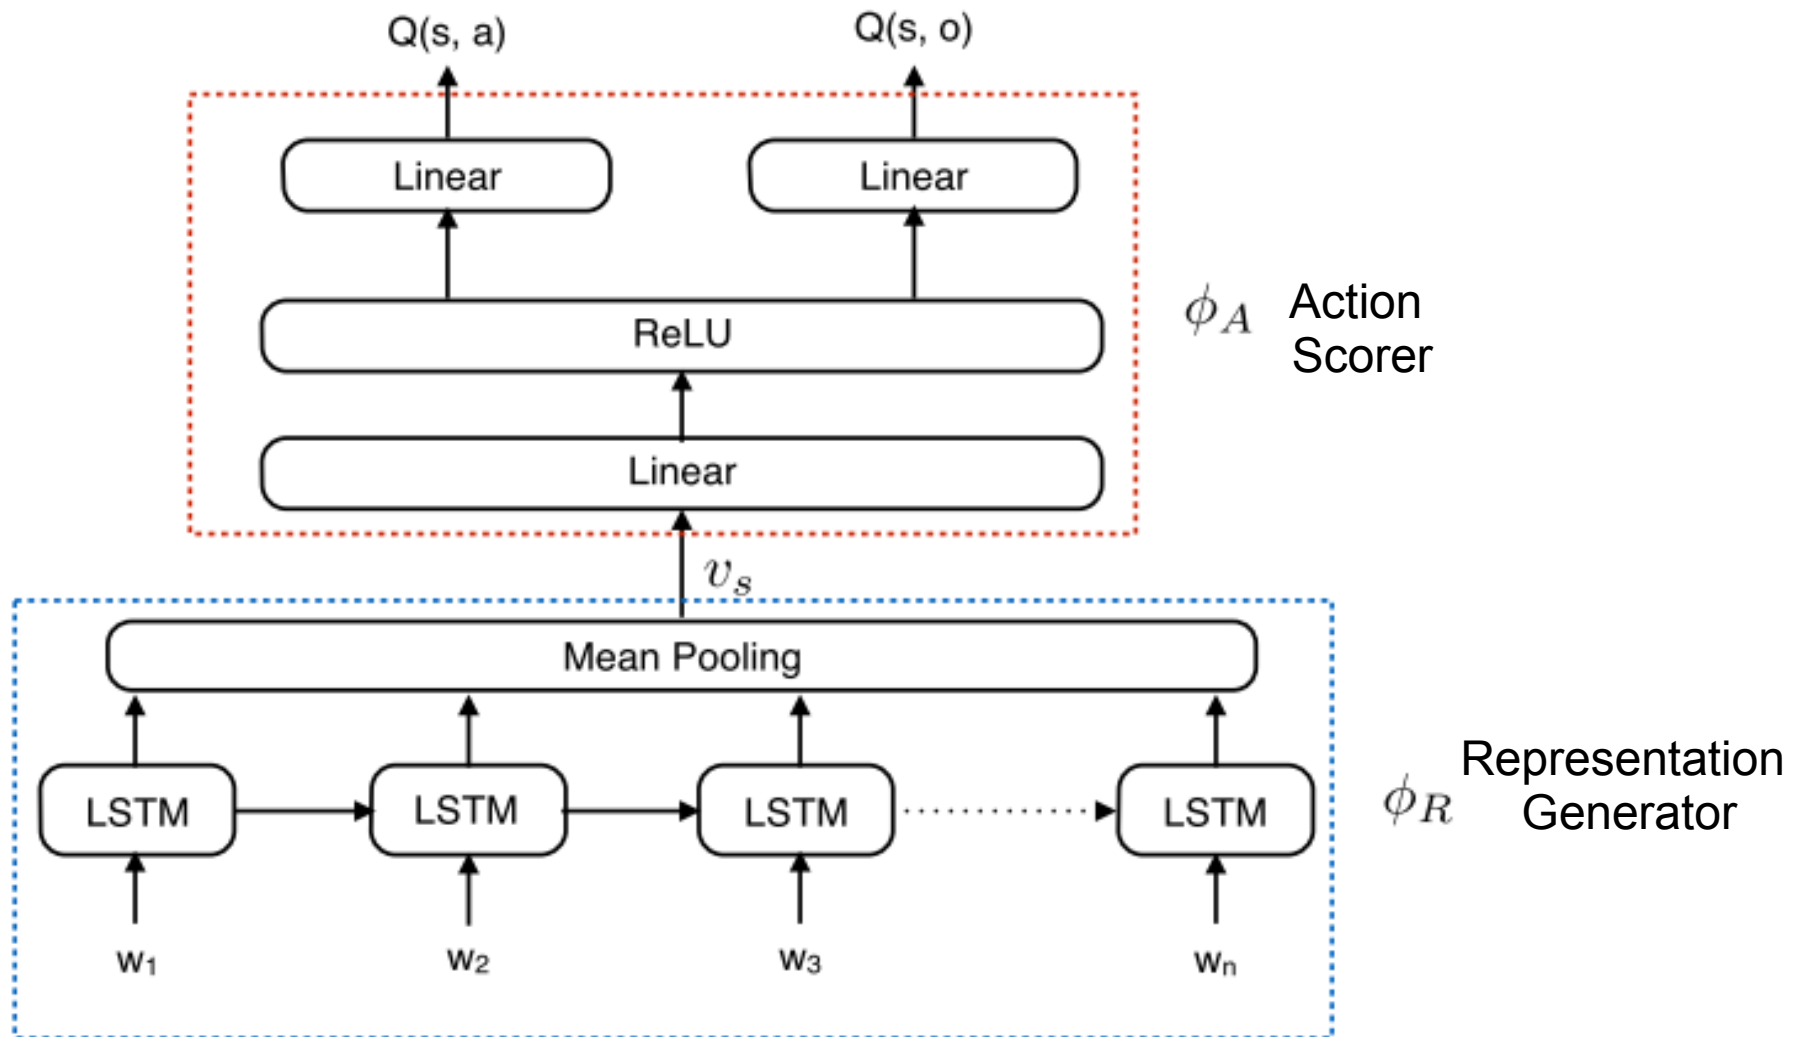


Long Short-Term Memory

- Write: $m_{in,t} = k_{out}C_{write}$
- Forget: $m_{out,t} = m_{in,t} + C_{forget}m_{out,t-1}$
- Read: $C_{out} = n_{out}C_{read}$



Deep Reinforcement Learning Framework

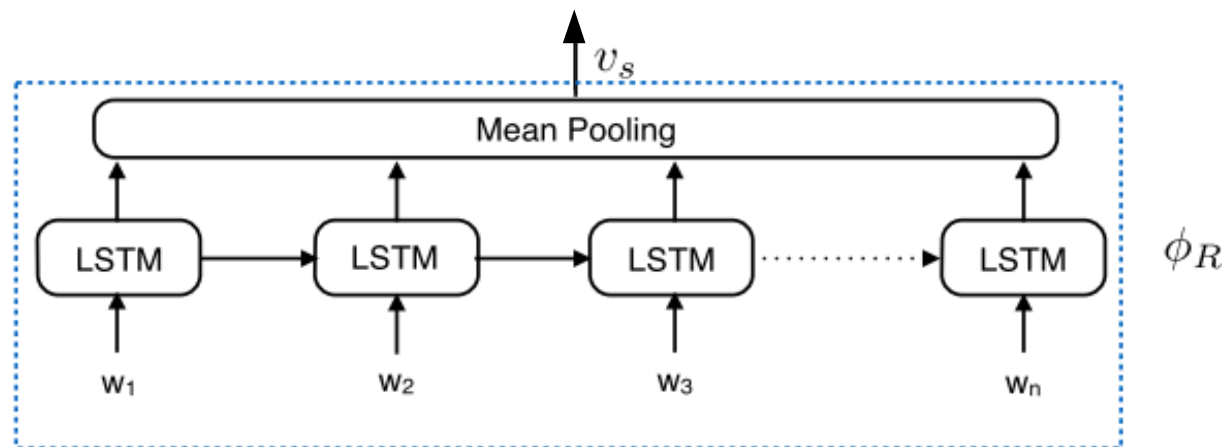


Game Representation

- Represented by the tuple $\langle H, A, T, R, \Psi \rangle$
- H : the set of all possible game states.
- $A = \{(a, o)\}$: all commands (action-object pairs).
- $T(h' \mid h, a, o)$: stochastic transition function between game states.
- $R(h, a, o)$: reward function.
- $\Psi : H \rightarrow S$: converts game state into a textual description S , since the game state is hidden.

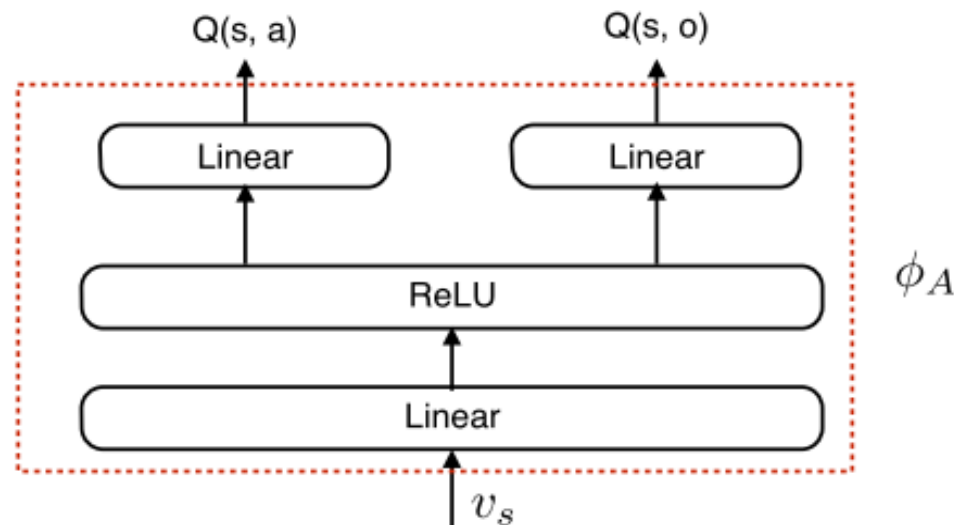
Representation Generator ϕ_R

- Reads displayed raw text w_1, w_2, \dots, w_n
- converts it to a vector representation v_s
- LSTM : $x_k = f(w_1, \dots, w_k)$
- mean pooling layer : $v_s = \frac{1}{n} \sum x_k$



Action Scorer ϕ_A

- multi-layered neural network
- given the current state representation v_s , produces scores for actions.
- only consider commands with one action and object : (a, o)



Parameter Learning

- learning the parameters θ_R, θ_A by using stochastic gradient descent with RMSprop
- keep track of the agent's previous experiences in a memory D
- sample a random transition $(\hat{s}, \hat{a}, s', r)$ from D , and updating the parameters.

$$L(\theta_i) = E_{\hat{s}, \hat{a}}[(r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(\hat{s}, \hat{a}; \theta_i))^2]$$

- prioritized sampling to the transition with reward $r > 0$

Parameter Learning

```
3: for episode = 1, M do
4:   Initialize game and get start state description  $s_1$ 
5:   for  $t = 1, T$  do
6:     Convert  $s_t$  (text) to representation  $v_{s_t}$  using  $\phi_R$ 
7:     if  $\text{random}() < \epsilon$  then
8:       Select a random action  $a_t$ 
9:     else
10:      Compute  $Q(s_t, a)$  for all actions using  $\phi_A(v_{s_t})$ 
11:      Select  $a_t = \operatorname{argmax} Q(s_t, a)$ 
12:      Execute action  $a_t$  and observe reward  $r_t$  and new state  $s_{t+1}$ 
13:      Set priority  $p_t = 1$  if  $r_t > 0$ , else  $p_t = 0$ 
14:      Store transition  $(s_t, a_t, r_t, s_{t+1}, p_t)$  in  $\mathcal{D}$ 
15:      Sample random mini batch of transitions  $(s_j, a_j, r_j, s_{j+1}, p_j)$  from  $\mathcal{D}$ ,
        with fraction  $\rho$  having  $p_j = 1$ 
16:      Set  $y_j = \begin{cases} r_j & \text{if } s_{j+1} \text{ is terminal} \\ r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta) & \text{if } s_{j+1} \text{ is non-terminal} \end{cases}$ 
17:      Perform gradient descent step on the loss  $\mathcal{L}(\theta) = (y_j - Q(s_j, a_j; \theta))^2$ 
```

Experiment & Result

- Two World:
 - Home world : hand-cracted simple testcase
 - Fantasy World : real MUD game testcase
- Evaluation:
 - the cumulative reward
 - the fraction of quests completed by the agent
- Baseline:
 - Random agent
 - Bag-of-words
 - bigram

Experiment & Result

- Home world :
 - Mimic a typical house with 4 rooms.
 - Each room contains a representative object.
 - Text : Description of the character's state and that the quest.

- Example:

quest:

Not you are sleepy
now but you are
hungry now.


command:

go to kitchen
eat apple


Experiment & Result

Stats	Home World	Fantasy World
Vocabulary size	84	1340
Avg. words / description	10.5	65.21
Max descriptions / room	3	100
# diff. quest descriptions	12	-
State transitions	Deterministic	Stochastic
# states (underlying)	16	≥ 56
Branching factor (# commands / state)	40	222


Experiment & Result



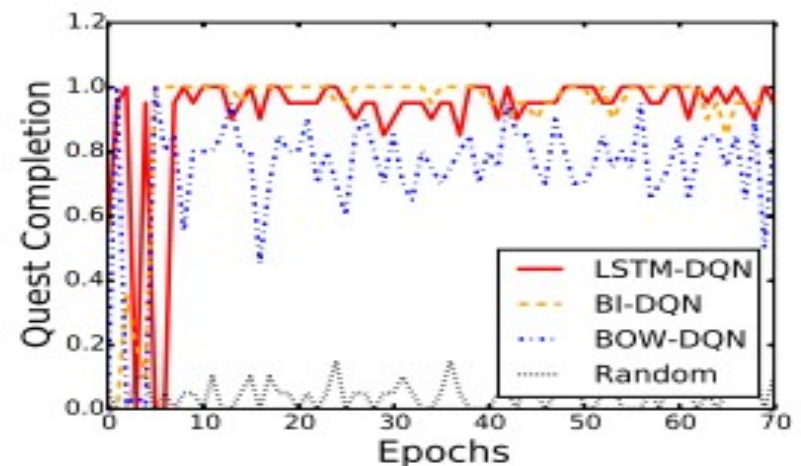
Reward (Home)



Quest completion (Home)




Reward (Fantasy)



Quest completion (Fantasy)

Experiment & Result



Conclusion

- a deep reinforcement learning framework
- jointly learn state representations and action policies using game rewards as feedback.
- mapping text descriptions into vectors that capture the semantics of the game states.